

Imports

```
In [25]: import os
import platform
import random
import kagglehub
import mlflow
import mlflow.pytorch
import optuna
import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import matplotlib.pyplot as plt
import tqdm as notebook_tqdm # Needed for tqdm in Jupyter Notebook (Certain
from PIL import Image
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms, models
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from torchvision.models import ResNet18_Weights
```

Hyperparameters Options

```
In [26]: DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
EXPERIMENT_NAME = "fire-smoke-detection-resnet-tuning"
SEED = 42
NUM_EPOCHS = 10
NUM_TRIALS = 1
BATCH_SIZE_OPTIONS = [16, 32, 64, 128]
LEARNING_RATE_OPTIONS = [1e-4, 1e-3, 1e-2]
WEIGHT_DECAY_OPTIONS = [1e-6, 1e-5, 1e-4]
EARLY_STOP_PATIENCE = 3
```

Download Dataset

```
In [27]: DATASET_PATH = kagglehub.dataset_download("sayedgamal99/smoke-fire-detection")
```

Data Augmentation Options

Input images are expected to be 224x224

```
In [28]: TRAIN_TRANSFORM = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
])
```

```
EVAL_TRANSFORM = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
])
```

Construct Custom Dataset

The original dataset is structure as such:

- [] = 'No Smoke and No Fire'
- 0 = 'Smoke Only'
- 1 = 'Fire and Smoke'

The custom dataset modifies this as such:

- 0 = 'No Smoke and No Fire'
- 1 = 'Smoke Only'
- 2 = 'Fire and Smoke'

```
In [29]: class CustomDataset(Dataset):
    def __init__(self, images_dir, labels_dir, transform=None):
        self.images_dir = images_dir
        self.labels_dir = labels_dir
        self.transform = transform
        self.image_files = sorted(os.listdir(images_dir))

    def __len__(self):
        return len(self.image_files)

    def __getitem__(self, idx):
        img_name = self.image_files[idx]
        img_path = os.path.join(self.images_dir, img_name)
        label_path = os.path.join(self.labels_dir, img_name.replace(".jpg",
                                                                    ".txt"))

        image = Image.open(img_path).convert("RGB")
        with open(label_path, "r") as f:
            label_content = f.read().strip()

        # 0: none, 1: smoke, 2: fire
        if not label_content:
            label = 0
        else:
            first_number = int(label_content.split()[0])
            label = 1 if first_number == 0 else 2

        if self.transform:
            image = self.transform(image)

        return image, label
```

Training Code

The training parameters are provided by the Optuna Trails

```
In [30]: def train_with_params(params: dict, train_dataset, val_dataset, test_dataset)
        """
        Train the model using the provided parameters and datasets.
        Returns the best validation loss.
        """

        batch_size = params["batch_size"]
        learning_rate = params["learning_rate"]
        weight_decay = params["weight_decay"]
        num_epochs = params["num_epochs"]
        early_stop_patience = params["early_stop_patience"]

        train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
        val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
        test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

        device = DEVICE
        model = models.resnet18(weights=ResNet18_Weights.DEFAULT)
        num_fters = model.fc.in_features
        model.fc = nn.Linear(num_fters, 3)
        model = model.to(device)
        criterion = nn.CrossEntropyLoss()
        optimizer = optim.Adam(model.parameters(), lr=learning_rate, weight_decay=weight_decay)
        scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min', patience=10)

        run_name = f"bs{batch_size}_lr{learning_rate:.0e}_wd{weight_decay:.0e}"
        with mlflow.start_run(nested=True, run_name=run_name):
            mlflow.log_param("learning_rate", learning_rate)
            mlflow.log_param("weight_decay", weight_decay)
            mlflow.log_param("batch_size", batch_size)
            mlflow.log_param("num_epochs", num_epochs)
            mlflow.log_param("early_stop_patience", early_stop_patience)
            mlflow.log_param("optimizer", optimizer.__class__.__name__)
            mlflow.log_param("scheduler", scheduler.__class__.__name__)
            mlflow.log_param("platform", platform.platform())
            mlflow.log_param("python_version", platform.python_version())
            print("Starting training...")
            print(f"Batch size: {batch_size}")
            print(f"Learning rate: {learning_rate}")
            print(f"Weight decay: {weight_decay}")
            print(f"Number of epochs: {num_epochs}")
            print(f"Early stop patience: {early_stop_patience}")
            print(f"Optimizer: {optimizer.__class__.__name__}")
            print(f"Scheduler: {scheduler.__class__.__name__}")

            best_val_loss = float('inf')
            best_model_state = None
            best_epoch = -1
            epochs_no_improve = 0
            train_losses = []
            val_losses = []

            for epoch in range(num_epochs):
```

```

model.train()
train_loss = 0.0
for batch_idx, (inputs, labels) in enumerate(train_loader, 1):
    inputs, labels = inputs.to(device), labels.to(device)
    optimizer.zero_grad()
    outputs = model(inputs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()
    train_loss += loss.item()
    train_losses.append(loss.item())
    mlflow.log_metric("batch_training_loss", loss.item(), step=epoch)

model.eval()
val_loss = 0.0
val_correct = 0
val_total = 0
with torch.no_grad():
    for val_inputs, val_labels in val_loader:
        val_inputs, val_labels = val_inputs.to(device), val_labels.to(device)
        val_outputs = model(val_inputs)
        v_loss = criterion(val_outputs, val_labels)
        val_loss += v_loss.item()
        _, val_predicted = torch.max(val_outputs, 1)
        val_correct += (val_predicted == val_labels).sum().item()
        val_total += val_labels.size(0)

avg_train_loss = train_loss / len(train_loader)
avg_val_loss = val_loss / len(val_loader)
val_losses.append(avg_val_loss)
val_accuracy = val_correct / val_total if val_total > 0 else 0
mlflow.log_metric("training_loss", avg_train_loss, step=epoch)
mlflow.log_metric("validation_loss", avg_val_loss, step=epoch)
mlflow.log_metric("validation_accuracy", val_accuracy, step=epoch)
print(f"Epoch {epoch+1}/{num_epochs}, Training Loss: {avg_train_loss}, Validation Loss: {avg_val_loss}, Validation Accuracy: {val_accuracy}")

checkpoint = {
    "epoch": epoch,
    "model_state_dict": model.state_dict(),
    "optimizer_state_dict": optimizer.state_dict(),
    "scheduler_state_dict": scheduler.state_dict(),
    "best_val_loss": best_val_loss,
}
checkpoint_path = f"checkpoint_epoch_{epoch+1}.pth"
torch.save(checkpoint, checkpoint_path)
mlflow.log_artifact(checkpoint_path)
os.remove(checkpoint_path)

if avg_val_loss < best_val_loss:
    epochs_no_improve = 0
    best_model_state = model.state_dict()
    best_val_loss = avg_val_loss
    best_epoch = epoch
else:
    epochs_no_improve += 1
    if epochs_no_improve >= early_stop_patience:

```

```

        break

    scheduler.step(avg_val_loss)
    best_val_accuracy = max(val_losses)
    mlflow.log_metric("best_val_accuracy", best_val_accuracy)
    mlflow.log_metric("learning_rate", optimizer.param_groups[0]['lr'])

    if best_model_state is not None:
        model.load_state_dict(best_model_state)
        best_checkpoint = {
            "epoch": best_epoch,
            "model_state_dict": model.state_dict(),
            "optimizer_state_dict": optimizer.state_dict(),
            "scheduler_state_dict": scheduler.state_dict(),
            "best_val_loss": best_val_loss,
        }
        torch.save(best_checkpoint, "best_model.pth")
        mlflow.log_artifact("best_model.pth")
        os.remove("best_model.pth")

    test_loss = 0.0
    test_correct = 0
    test_total = 0
    with torch.no_grad():
        for test_inputs, test_labels in test_loader:
            test_inputs, test_labels = test_inputs.to(device), test_labels.to(device)
            outputs = model(test_inputs)
            loss = criterion(outputs, test_labels)
            test_loss += loss.item()
            _, predicted = torch.max(outputs, 1)
            test_correct += (predicted == test_labels).sum().item()
            test_total += test_labels.size(0)
    avg_test_loss = test_loss / len(test_loader)
    test_accuracy = test_correct / test_total if test_total > 0 else 0
    mlflow.log_metric("test_loss", avg_test_loss)
    mlflow.log_metric("test_accuracy", test_accuracy)
    print(f"Test Loss: {avg_test_loss:.4f}, Test Accuracy: {test_accuracy:.4f}")

    # Plot the confusion matrix
    y_true, y_pred = [], []
    with torch.no_grad():
        for test_inputs, test_labels in test_loader:
            test_inputs, test_labels = test_inputs.to(device), test_labels.to(device)
            outputs = model(test_inputs)
            _, predicted = torch.max(outputs, 1)
            y_true.extend(test_labels.cpu().numpy())
            y_pred.extend(predicted.cpu().numpy())
    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm)
    disp.plot(cmap=plt.cm.Blues)
    plt.title("Test Confusion Matrix")
    plt.tight_layout()
    plt.savefig("test_confusion_matrix.png")
    plt.show()
    mlflow.log_artifact("test_confusion_matrix.png")
    plt.close()

```

```

os.remove("test_confusion_matrix.png")

# Calculate average training loss per epoch
num_batches_per_epoch = len(train_loader)
train_loss_per_epoch = [
    np.mean(train_losses[i * num_batches_per_epoch : (i + 1) * num_batches_per_epoch])
    for i in range(len(val_losses))
]

# Plot training loss per epoch
plt.figure()
plt.plot(range(1, len(train_loss_per_epoch) + 1), train_loss_per_epoch)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Training Loss per Epoch")
plt.legend()
plt.tight_layout()
plt.savefig("training_loss_per_epoch.png")
plt.show()
mlflow.log_artifact("training_loss_per_epoch.png")
plt.close()
os.remove("training_loss_per_epoch.png")

# Plot validation loss per epoch
plt.figure()
plt.plot(range(1, len(val_losses) + 1), val_losses, marker="o", color="red")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Validation Loss per Epoch")
plt.legend()
plt.tight_layout()
plt.savefig("validation_loss_per_epoch.png")
plt.show()
mlflow.log_artifact("validation_loss_per_epoch.png")
plt.close()
os.remove("validation_loss_per_epoch.png")

return best_val_loss

```

Experiment and Trails Set-Up

```

In [31]: def objective(trial, train_dataset, val_dataset, test_dataset):
    params = {
        "batch_size": trial.suggest_categorical("batch_size", BATCH_SIZE_OPTIONS),
        "learning_rate": trial.suggest_float("learning_rate", LEARNING_RATE_MIN, LEARNING_RATE_MAX),
        "weight_decay": trial.suggest_float("weight_decay", WEIGHT_DECAY_MIN, WEIGHT_DECAY_MAX),
        "num_epochs": NUM_EPOCHS,
        "early_stop_patience": EARLY_STOP_PATIENCE,
    }
    return train_with_params(params, train_dataset, val_dataset, test_dataset)

def start_experiment():
    # Set seed for reproducibility
    random.seed(SEED)

```

```

np.random.seed(SEED)
torch.manual_seed(SEED)
if torch.cuda.is_available():
    torch.cuda.manual_seed_all(SEED)

# Load training dataset
train_dataset = CustomDataset(
    images_dir= os.path.join(DATASET_PATH, "data/train/images"),
    labels_dir= os.path.join(DATASET_PATH, "data/train/labels"),
    transform=TRAIN_TRANSFORM
)

# Load validation dataset
val_dataset = CustomDataset(
    images_dir= os.path.join(DATASET_PATH, "data/val/images"),
    labels_dir= os.path.join(DATASET_PATH, "data/val/labels"),
    transform=EVAL_TRANSFORM
)

# Load test dataset
test_dataset = CustomDataset(
    images_dir= os.path.join(DATASET_PATH, "data/test/images"),
    labels_dir= os.path.join(DATASET_PATH, "data/test/labels"),
    transform=EVAL_TRANSFORM
)

# Create study
study = optuna.create_study(direction="minimize", study_name=EXPERIMENT_NAME)
mlflow.set_experiment(EXPERIMENT_NAME)
study.optimize(
    lambda trial: objective(trial, train_dataset, val_dataset, test_dataset),
    n_trials=NUM_TRIALS
)

# Print best trial
print("Best trial:")
print(f" Value (best validation loss): {study.best_trial.value}")
print(" Params: ")
for key, value in study.best_trial.params.items():
    print(f"    {key}: {value}")

# Log best trial info with MLflow
mlflow.log_metric("best_val_loss", study.best_trial.value)
for key, value in study.best_trial.params.items():
    mlflow.log_param(f"best_{key}", value)



```


Start the Experiment

In [32]: `start_experiment()`

[I 2025-04-28 20:37:20,569] A new study created in memory with name: fire-smoke-detection-resnet-tuning

```
Starting training...
Batch size: 128
Learning rate: 0.004840004649523287
Weight decay: 3.0846377146227715e-06
Number of epochs: 10
Early stop patience: 3
Optimizer: Adam
Scheduler: ReduceLRonPlateau
Epoch 1/10, Training Loss: 0.9562, Validation Loss: 1.0056, Validation Accuracy: 0.5115
Epoch 2/10, Training Loss: 0.7917, Validation Loss: 1.0393, Validation Accuracy: 0.5073
Epoch 3/10, Training Loss: 0.6925, Validation Loss: 0.7461, Validation Accuracy: 0.6592
Epoch 4/10, Training Loss: 0.6646, Validation Loss: 0.8331, Validation Accuracy: 0.6299
Epoch 5/10, Training Loss: 0.5825, Validation Loss: 0.6321, Validation Accuracy: 0.7128
Epoch 6/10, Training Loss: 0.5476, Validation Loss: 0.6347, Validation Accuracy: 0.7131
Epoch 7/10, Training Loss: 0.5006, Validation Loss: 0.6384, Validation Accuracy: 0.6957
Epoch 8/10, Training Loss: 0.4027, Validation Loss: 0.4289, Validation Accuracy: 0.8319
Epoch 9/10, Training Loss: 0.3574, Validation Loss: 0.4156, Validation Accuracy: 0.8309
Epoch 10/10, Training Loss: 0.3410, Validation Loss: 0.4112, Validation Accuracy: 0.8361
Test Loss: 0.3974, Test Accuracy: 0.8453
```



[I 2025-04-28 20:58:08,042] Trial 0 finished with value: 0.41120924927294256 and parameters: {'batch_size': 128, 'learning_rate': 0.004840004649523287, 'weight_decay': 3.0846377146227715e-06}. Best is trial 0 with value: 0.41120924927294256.

Best trial:

Value (best validation loss): 0.41120924927294256

Params:

batch_size: 128

learning_rate: 0.004840004649523287

weight_decay: 3.0846377146227715e-06